

COMP9417: Homework Set #2

z5113817

University of New South Wales — July 18, 2021

All code for this homework set is available [here](#).

Question 1


a

The possible values for both y_i and \tilde{y}_i are binary. Even though they have different values ($y_i \in \{0, 1\}$ and $\tilde{y}_i \in \{-1, 1\}$), the objective of each logistic regression implementation is to divide the dataset into 2 classifications. Because of this, the actual value that each classification has will not affect the parameters that the regression is attempting to optimise ($(\hat{\beta}_0, \hat{\beta})$ and (\hat{w}, \hat{c})). Therefore the solutions for the parameters being minimised by each regression will be the same.

C is a hyper-parameter that adjusts the sensitivity that the model has to its coefficients. Compared with the standard LASSO parameter λ , C is a multiple of the Loss function whereas λ is a multiple of the Penalty.

b

Boxplot of testing accuracy for each value of C :



The value of C returning best results: **0.18794747474747472**

The testing accuracy of this model: **76%**

c

From GridSearchCV:

The value of C returning best results: **0.0122191919191918**

The testing accuracy of this model: **75.2%**

In our answer for b , we determined the "best" value of C as the value which corresponded to the average lowest *log-loss* value across all folds. The value from the *GridSearchCV* are different because, by default, it determines the "best" value of C as the one which corresponds to the average highest *score* across all folds ¹.

We can modify the *GridSearchCV* class by providing our own metric for *scoring*. The following code is a scorer that uses the smallest *log-loss* value as its scoring metric.

```
scoring = make_scorer(  
    log_loss,                # The sklearn implementation of log_loss  
    greater_is_better=False, # A smaller log_loss is a better value  
    needs_proba=True        # Calculating log_loss needs the probability predictions  
)
```

This yields the following results:

```
grid_lr.best_estimator_ : LogisticRegression(C=0.181887878787877, penalty='l1', solver='liblinear')  
grid_lr.best_score_ (lowest log-loss): -0.5374067706086373
```

This value of *C* matches my value in b .

¹See *scoring* parameter in documentation. If the estimator provided exposes a *score* method and a value for *scoring* is not provided, then *score* is used to determine the "best" value of C

Code For Questions

q1b

```
21 def q1b():
22     # Import data and take first 500 rows
23     data = import_data(Q1_DATA_DIR)
24     data = data.head(500)
25
26     # Create folds
27     folds = []
28     fold_size = 50
29     for i in range(0,10):
30         folds.append(data[i*fold_size:i*fold_size+fold_size])
31
32     # Create grid of 100 C values
33     C_grid = linspace(0.0001, 0.6, 100)
34
35     # Iterate over C values and train Logistic Model
36     scores = []
37     for C in C_grid.tolist():
38         # Create model
39         model = LogisticRegression(penalty="l1", solver="liblinear", C=C)
40
41         # Iterate over each fold to train and test on each
42         fold_scores = []
43         for i in range(0, len(folds)):
44             # Set up training and test data
45             folds_train = folds.copy() # Create a temp so that we don't mutate original 'folds'
46             test_df = folds_train.pop(i)
47             train_df = concat(folds_train)
48
49             # Fit model to data
50             clf = model.fit(train_df.drop("Y", axis=1), train_df["Y"])
51
52             # Score data
53             clf_probs = clf.predict_proba(test_df.drop("Y", axis=1))
54             fold_score = log_loss(test_df["Y"], clf_probs)
55             fold_scores.append(fold_score)
56
57         scores.append(fold_scores)
```

```

58
59 # Save boxplot
60 fig, ax = plt.subplots()
61 ax.boxplot(scores)
62 ax.set_xticklabels([""]) You, 21 hours ago • Finish q1c
63 ax.set_title("Logistic Regression Log-Loss For Various C Values")
64 ax.set_xlabel("C Values")
65 ax.set_ylabel("Log-Loss")
66 plt.savefig("./outputs/q1b_boxplot.png")
67
68 # Record the value of C that returns the "best" CV score
69 #
70 # "Best" is taken to be the lowest average log-loss
71 averages = list(map(lambda x : mean(x), scores))
72
73 # Get index of lowest average
74 i = averages.index(min(averages))
75
76 # Map this to a C value
77 C_best = C_grid[i]
78 print(C_best)
79
80 # Retrain this model and report its accuracy
81 model = LogisticRegression(penalty="l1", solver="liblinear", C=C_best)
82
83 # Let's test on the first fold and train on the remainder
84 test_df = folds.pop(0)
85 train_df = concat(folds)
86 clf = model.fit(train_df.drop("Y", axis=1), train_df["Y"])
87
88 print(clf.score(test_df.drop("Y", axis=1), test_df["Y"]))

```

q1c

```

91 def q1c():
92     # Import data
93     data = import_data(Q1_DATA_DIR)
94     data = data.head(500)
95     Xtrain = data.drop("Y", axis=1)
96     Ytrain = data["Y"]
97
98     # Create grid of 100 C values
99     C_grid = linspace(0.0001, 0.6, 100)
100     param_grid = { "C": C_grid }
101
102     # Set our own scoring metric
103     scoring = make_scorer(
104         log_loss,
105         greater_is_better=False,
106         needs_proba=True
107     )
108
109     # Assignment code
110     grid_lr = GridSearchCV(
111         estimator=LogisticRegression(penalty='l1', solver='liblinear'),
112         cv=10,
113         param_grid=param_grid,
114         scoring=scoring
115     )
116     grid_lr.fit(Xtrain, Ytrain)
117     print(grid_lr.best_estimator_)
118     print(grid_lr.best_score_)

```